

Ved Nigam Introduction to ML in Python with NumPy, Pandas, Seaborn, an SKLearn

```
In [ ]: import pandas as pd
import numpy as np

# 1. Reading in data
# a.
df = pd.read_csv('Auto.csv')

# b.
print(df.head())

#c.
print(df.shape)
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	
0	18.0	8	307.0	130	3504	12.0	70.0	\
1	15.0	8	350.0	165	3693	11.5	70.0	
2	18.0	8	318.0	150	3436	11.0	70.0	
3	16.0	8	304.0	150	3433	12.0	70.0	
4	17.0	8	302.0	140	3449	NaN	70.0	

	origin	name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino

(392, 9)

```
In [ ]: # 2. Data Exploration
# a. b.
df_mpg = df[['mpg']]
print(df_mpg.describe())
# average mpg is 23.445918 and the range for mpg in the data is 37

df_weight = df[['weight']]
print(df_weight.describe())
# average weight is 2977.584184 and the range for weight is 3527

df_year = df[['year']]
print(df_year.describe())
# average year is 76.010256 and the range is 12
```

```
      mpg
count  392.000000
mean   23.445918
std     7.805007
min     9.000000
25%    17.000000
50%    22.750000
75%    29.000000
max    46.600000

      weight
count  392.000000
mean  2977.584184
std   849.402560
min  1613.000000
25%  2225.250000
50%  2803.500000
75%  3614.750000
max  5140.000000

      year
count  390.000000
mean   76.010256
std     3.668093
min    70.000000
25%    73.000000
50%    76.000000
75%    79.000000
max    82.000000
```

```
In [ ]: # 3.
        # a.
        print(df.dtypes)

        # b.
        df.cylinders = df.cylinders.astype('category').cat.codes

        # c.
        df['origin'] = pd.Categorical(df.origin)

        # d.
        print(df.dtypes)
```

```

mpg          float64
cylinders     int64
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year          float64
origin        int64
name          object
dtype: object
mpg          float64
cylinders     int8
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year          float64
origin        category
name          object
dtype: object

```

```

In [ ]: # 4.
        # a.
        df = df.dropna()

        # b.
        print(df.shape)

```

(389, 9)

```

In [ ]: # 5.
        # a.
        df['mpg_high'] = np.where(df.mpg > np.mean(df.mpg), 1, 0)
        del df['mpg']
        del df['name']

        # b.
        df.head

```

```
Out[ ]: <bound method NDFrame.head of
ght acceleration year origin cylinders displacement horsepower wei
0          4      307.0         130      3504          12.0  70.0         1
\
1          4      350.0         165      3693          11.5  70.0         1
2          4      318.0         150      3436          11.0  70.0         1
3          4      304.0         150      3433          12.0  70.0         1
6          4      454.0         220      4354           9.0  70.0         1
..      ...      ...      ...      ...      ...      ...
387        1      140.0          86      2790          15.6  82.0         1
388        1       97.0          52      2130          24.6  82.0         2
389        1      135.0          84      2295          11.6  82.0         1
390        1      120.0          79      2625          18.6  82.0         1
391        1      119.0          82      2720          19.4  82.0         1
```

```
mpg_high
0          0
1          0
2          0
3          0
6          0
..      ...
387        1
388        1
389        1
390        1
391        1
```

[389 rows x 8 columns]>


```
In [ ]: import seaborn as sb


# 6.
# a.
sb.catplot(y = 'mpg_high', data = df)
# Data is either 0 or 1

# b.
sb.relplot(x = 'horsepower', y = 'weight', hue = 'mpg_high', data = df)
# Higher horsepower cars are less fuel efficient

# c.
sb.boxplot(x = 'mpg_high', y = 'weight', data = df)
# The fuel efficient cars are lighter
```

```
Out[ ]: <Axes: xlabel='mpg_high', ylabel='weight'>
```





```
In [ ]: from sklearn.model_selection import train_test_split

X = df.iloc[:, 0:6]
y = df.iloc[:, 7]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print('train size:', X_train.shape)
print('test size:', X_test.shape)

print(X_train)
```

```

train size: (311, 6)
test size: (78, 6)
   cylinders  displacement  horsepower  weight  acceleration  year
184         1         101.0           83   2202         15.3   76.0
355         3         145.0           76   3160         19.6   81.0
57          1          97.5           80   2126         17.0   72.0
170         1          90.0           71   2223         16.5   75.0
210         4         350.0          180   4380         12.1   76.0
..         ...         ...         ...         ...         ...
207         1         120.0           88   3270         21.9   76.0
56          1         113.0           95   2278         15.5   72.0
297         1         141.0           71   3190         24.8   79.0
214         1          98.0           68   2045         18.5   77.0
306         1         151.0           90   2556         13.2   79.0

```

[311 rows x 6 columns]

```

In [ ]: # 8
from sklearn.linear_model import LogisticRegression

clf = LogisticRegression()
clf.fit(X_train, y_train)
print(clf.score(X_train, y_train))

pred = clf.predict(X_test)

from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
print('accuracy score: ', accuracy_score(y_test, pred))
print('precision score: ', precision_score(y_test, pred))
print('recall score: ', recall_score(y_test, pred))
print('f1 score: ', f1_score(y_test, pred))

```

0.9035369774919614

accuracy score: 0.8589743589743589

precision score: 0.7297297297297297

recall score: 0.9642857142857143

f1 score: 0.8307692307692307

/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```

In [ ]: # 9.
from sklearn.tree import DecisionTreeClassifier

clf_tree = DecisionTreeClassifier()
clf_tree.fit(X_train, y_train)

pred_tree = clf.predict(X_test)

```

```

from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

print('accuracy score: ', accuracy_score(y_test, pred_tree))
print('precision score: ', precision_score(y_test, pred_tree))
print('recall score: ', recall_score(y_test, pred_tree))
print('f1 score: ', f1_score(y_test, pred_tree))

```

```

accuracy score: 0.8589743589743589
precision score: 0.7297297297297297
recall score: 0.9642857142857143
f1 score: 0.8307692307692307

```

```

In [ ]: # 10.
from sklearn import preprocessing

scaler = preprocessing.StandardScaler().fit(X_train)

X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)

from sklearn.neural_network import MLPClassifier

clf = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(5, 2), max_iter=500,
clf.fit(X_train_scaled, y_train)

pred = clf.predict(X_test_scaled)

print('accuracy = ', accuracy_score(y_test, pred))

from sklearn.metrics import classification_report
print(classification_report(y_test, pred))

clf_sgd = MLPClassifier(solver='sgd', hidden_layer_sizes=(5, 2), max_iter=500,
clf_sgd.fit(X_train_scaled, y_train)

pred_sgd = clf_sgd.predict(X_test_scaled)

print('accuracy = ', accuracy_score(y_test, pred_sgd))

print(classification_report(y_test, pred_sgd))

```



```

accuracy = 0.8846153846153846
          precision    recall  f1-score   support

         0           0.94      0.88      0.91         50
         1           0.81      0.89      0.85         28

   accuracy                0.88         78
  macro avg           0.87      0.89      0.88         78
 weighted avg           0.89      0.88      0.89         78

accuracy = 0.8846153846153846
          precision    recall  f1-score   support

         0           0.94      0.88      0.91         50
         1           0.81      0.89      0.85         28

   accuracy                0.88         78
  macro avg           0.87      0.89      0.88         78
 weighted avg           0.89      0.88      0.89         78

```

```

/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-packages/sklearn/neural_network/_multilayer_perceptron.py:686: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (500) reached and the optimization hasn't converged yet.
  warnings.warn(

```

11.

The 2 neural networks got the same response. In terms of efficiency in algorithms, the logistic regression using Sci kit learn was the fastest, next to decision trees and the neural networks. Maybe with a larger dataset, the difference will be more prevalent. Decision Trees are known to be a greedy algorithm because of how they work by splitting at every single node. With datasets larger than normal, this can get highly inefficient. I was expecting the Neural Networks to be the fastest, maybe I did something wrong. I wish I had more time to investigate, but it is definitely something I will be observant of as I continue to experiment with algorithms vs Neural Networks in later projects.

I don't see a significant difference between R and Python so far because it is all just knowing a bunch of keywords for me. Having coded in Python, this is not coding in Python, but a good example of syntactically powerful the language us. At the stage we're at of ML in Python, I would prefer R because I like to be able to see the data/datatypes/workspace in R Studios. Although that is an IDE difference, that is really the only difference I've noticed. There is more documentation online for Python, which is definitely a plus point.